



中山大學
SUN YAT-SEN UNIVERSITY

Lecture 7

MapReduce Algorithms

SE-808 Cloud Application Development (supported by Google)

<http://my.ss.sysu.edu.cn/courses/cloud/>

School of Software, Sun Yat-sen University

Outline

- Basic MapReduce Algorithms
 - Sort/Search
- MapReduce algorithm design
 - Managing dependencies
 - Coordinating mappers and reducers
- Case study #1: inverted index
- Case study #2: pairwise similarity comparison

Part of the slides are adapted from Jimmy Lin's cloud course slides: <http://www.umiacs.umd.edu/~jimmylin/cloud-2010-Spring/index.html>

MapReduce Jobs

- Tend to be very short, code-wise
 - IdentityReducer is very common
- “Utility” jobs can be composed
- Represent a data flow, more so than a procedure

Sort: Inputs


- A set of files, one value per line.
- Mapper key is file name, line number.
- Mapper value is the contents of the line.

Sort Algorithm

- Takes advantage of reducer properties:
 - (key, value) pairs are processed in order by key;
 - reducers are themselves ordered •
- Mapper
 - Identity function for value $(k, v) \rightarrow (v, _)$
- Reducer
 - Identity function $(k', _) \rightarrow (k', _)$

Sort: The Trick

- (key, value) pairs from mappers are sent to a particular reducer based on hash(key)
- Must pick the hash function for your data such that
 - $k < k' \Rightarrow \text{hash}(k) < \text{hash}(k')$



Final Thoughts on Sort

- Used as a test of Hadoop's raw speed
 - Essentially "IO drag race"
 - Highlights utility of GFS
-
- Yahoo: Hadoop Sorts a Petabyte in 16.25 Hours and a Terabyte in 62 Seconds, 3800 nodes, May 2009.

Search: Inputs

- A set of files containing lines of text
 - A search pattern to find
 - Mapper key is file name, line number
 - Mapper value is the contents of the line
- Search pattern sent as special parameter

Search Algorithm

- Mapper:
Given (filename, some text) and “pattern”, if
“text” matches “pattern” output (filename, _)
- Reducer:
Identity function

Exercise 1: Graph reversal

- Given a directed graph as an adjacency list:

src1: dest11, dest12, ...
src2: dest21, dest22, ...
- Construct the graph in which all the links are reversed

Exercise 2: Frequent Pairs

- Given a large set of market baskets, find all frequent pairs
 - Frequent pairs are item pairs with counts greater than a given threshold

MapReduce Algorithm Design

Managing Dependencies

- Remember: Mappers run in isolation
 - You have no idea in what order the mappers run
 - You have no idea on what node the mappers run
 - You have no idea when each mapper finishes
- Tools for synchronization:
 - Ability to hold state in reducer across multiple key-value pairs
 - Sorting function for keys
 - Partitioner
 - Cleverly-constructed data structures

Motivating Example

- Term co-occurrence matrix for a text collection
 - $M = N \times N$ matrix (N = vocabulary size)
 - M_{ij} : number of times i and j co-occur in some context (for concreteness, let's say context = sentence)
- Why?
 - Distributional profiles as a way of measuring semantic distance
 - Semantic distance useful for many language processing tasks

“You shall know a word by the company it keeps” (Firth, 1957)

e.g., Mohammad and Hirst (EMNLP, 2006)

MapReduce: Large Counting Problems

- Term co-occurrence matrix for a text collection
= specific instance of a large counting problem
 - A large event space (number of terms)
 - A large number of events (the collection itself)
 - Goal: keep track of interesting statistics about the events
- Basic approach
 - Mappers generate partial counts
 - Reducers aggregate partial counts

How do we aggregate partial counts efficiently?

First Try: “Pairs”

- Each mapper takes a sentence:
 - Generate all co-occurring term pairs
 - For all pairs, emit $(a, b) \rightarrow \text{count}$
- Reducers sums up counts associated with these pairs
- Use combiners!

Note: in all these slides, a key-value pair denoted as $k \rightarrow v$

“Pairs” Analysis

- Advantages
 - Easy to implement, easy to understand
- Disadvantages
 - Lots of pairs to sort and shuffle around (upper bound?)

Another Try: “Stripes”

- Idea: group together pairs into an associative array

(a, b) \rightarrow 1

(a, c) \rightarrow 2

(a, d) \rightarrow 5

(a, e) \rightarrow 3

(a, f) \rightarrow 2


$a \rightarrow \{ b: 1, c: 2, d: 5, e: 3, f: 2 \}$

- Each mapper takes a sentence:
 - Generate all co-occurring term pairs
 - For each term, emit $a \rightarrow \{ b: \text{count}_b, c: \text{count}_c, d: \text{count}_d \dots \}$
- Reducers perform element-wise sum of associative arrays

$$\begin{array}{r}
 a \rightarrow \{ b: 1, \quad d: 5, e: 3 \} \\
 + \quad a \rightarrow \{ b: 1, c: 2, d: 2, \quad f: 2 \} \\
 \hline
 a \rightarrow \{ b: 2, c: 2, d: 7, e: 3, f: 2 \}
 \end{array}$$

“Stripes” Analysis

- Advantages
 - Far less sorting and shuffling of key-value pairs
 - Can make better use of combiners
- Disadvantages
 - More difficult to implement
 - Underlying object is more heavyweight
 - Fundamental limitation in terms of size of event space

Efficiency comparison of approaches to computing word co-occurrence matrices

Cluster size: 38 cores

Data Source: Associated Press Worldstream (APW) of the English Gigaword Corpus (v3), which contains 2.27 million documents (1.8 GB compressed, 5.7 GB uncompressed)

Conditional Probabilities

- How do we compute conditional probabilities from counts?

$$P(B | A) = \frac{\text{count}(A, B)}{\text{count}(A)} = \frac{\text{count}(A, B)}{\sum_{B'} \text{count}(A, B')}$$

- Why do we want to do this?
- How do we do this with MapReduce?

P(B|A): “Pairs”

$(a, *) \rightarrow 32$

Reducer holds this value in memory

$(a, b_1) \rightarrow 3$

$(a, b_2) \rightarrow 12$

$(a, b_3) \rightarrow 7$

$(a, b_4) \rightarrow 1$

...



$(a, b_1) \rightarrow 3 / 32$

$(a, b_2) \rightarrow 12 / 32$

$(a, b_3) \rightarrow 7 / 32$

$(a, b_4) \rightarrow 1 / 32$

...

- For this to work:
 - Must emit extra $(a, *)$ for every b_n in mapper
 - Must make sure all a 's get sent to same reducer (use partitioner)
 - Must make sure $(a, *)$ comes first (define sort order)
 - Must hold state in reducer across different key-value pairs

P(B|A): “Stripes”

$a \rightarrow \{b_1:3, b_2:12, b_3:7, b_4:1, \dots\}$

- Easy!
 - One pass to compute $(a, *)$
 - Another pass to directly compute $P(B|A)$

Synchronization in Hadoop

- Approach 1: turn synchronization into an ordering problem
 - Sort keys into correct order of computation
 - Partition key space so that each reducer gets the appropriate set of partial results
 - Hold state in reducer across multiple key-value pairs to perform computation
 - Illustrated by the “pairs” approach
- Approach 2: construct data structures that “bring the pieces together”
 - Each reducer receives all the data it needs to complete the computation
 - Illustrated by the “stripes” approach

Issues and Tradeoffs

- Number of key-value pairs
 - Object creation overhead
 - Time for sorting and shuffling pairs across the network
- Size of each key-value pair
 - De/serialization overhead
- Combiners make a big difference!
 - RAM vs. disk and network
 - Arrange data to maximize opportunities to aggregate partial results

Questions?

Case study #1: Inverted Index

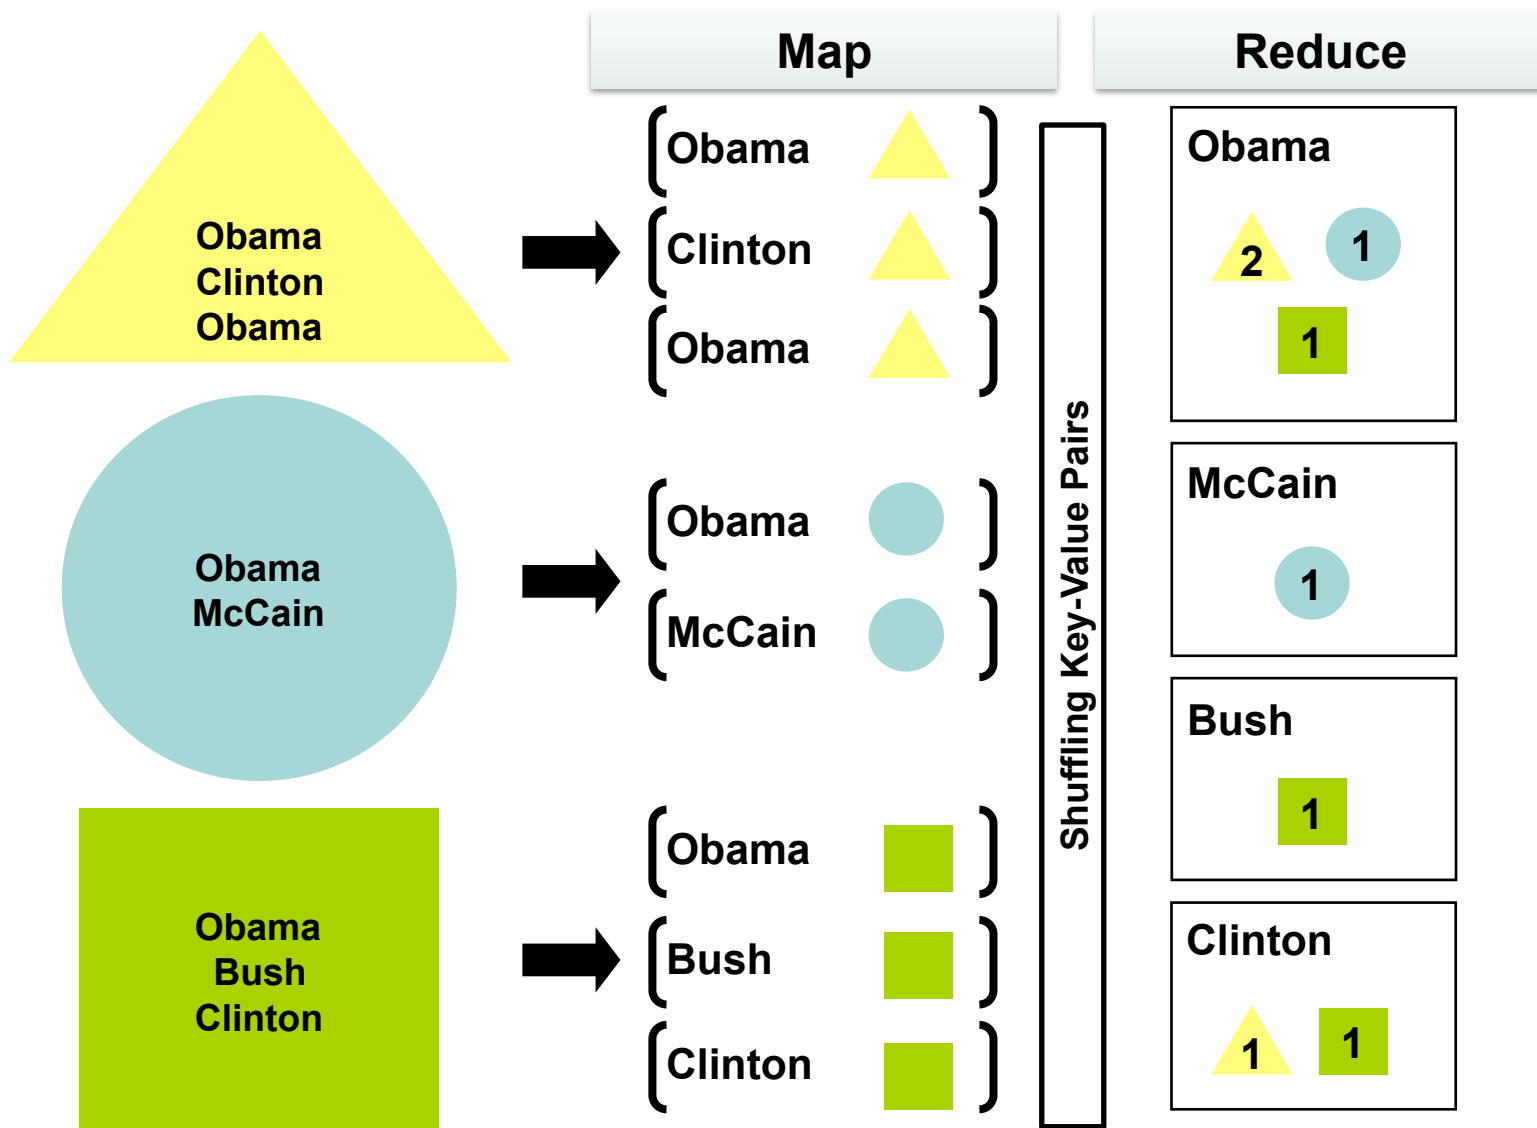
Inverted Index

Term	Doc 1	Doc 2	Doc 3	Doc 4	Doc 5	Doc 6	Doc 7	Doc 8
aid	0	0	0	1	0	0	0	1
all	0	1	0	1	0	1	0	0
back	1	0	1	0	0	0	1	0
brown	1	0	1	0	1	0	1	0
come	0	1	0	1	0	1	0	1
dog	0	0	1	0	1	0	0	0
fox	0	0	1	0	1	0	1	0
good	0	1	0	1	0	1	0	1
jump	0	0	1	0	0	0	0	0
lazy	1	0	1	0	1	0	1	0
men	0	1	0	1	0	0	0	1
now	0	1	0	0	0	1	0	1
over	1	0	1	0	1	0	1	1
party	0	0	0	0	0	1	0	1
quick	1	0	1	0	0	0	0	0
their	1	0	0	0	1	0	1	0



Term	Postings							
aid	→ 4	→ 8						
all	→ 2	→ 4	→ 6					
back	→ 1	→ 3	→ 7					
brown	→ 1	→ 3	→ 5	→ 7				
come	→ 2	→ 4	→ 6	→ 8				
dog	→ 3	→ 5						
fox	→ 3	→ 5	→ 7					
good	→ 2	→ 4	→ 6	→ 8				
jump	→ 3							
lazy	→ 1	→ 3	→ 5	→ 7				
men	→ 2	→ 4	→ 8					
now	→ 2	→ 6	→ 8					
over	→ 1	→ 3	→ 5	→ 7	→ 8			
party	→ 6	→ 8						
quick	→ 1	→ 3						
their	→ 1	→ 5	→ 7					


Map Reduce Algorithm: Inverted Index



Questions?

Case study #2: pairwise similarity comparison

Pairwise Document Similarity



- Applications:
 - Clustering
 - Cross-document coreference resolution
 - “more-like-that” queries

Problem Description

- Consider similarity functions of the form:

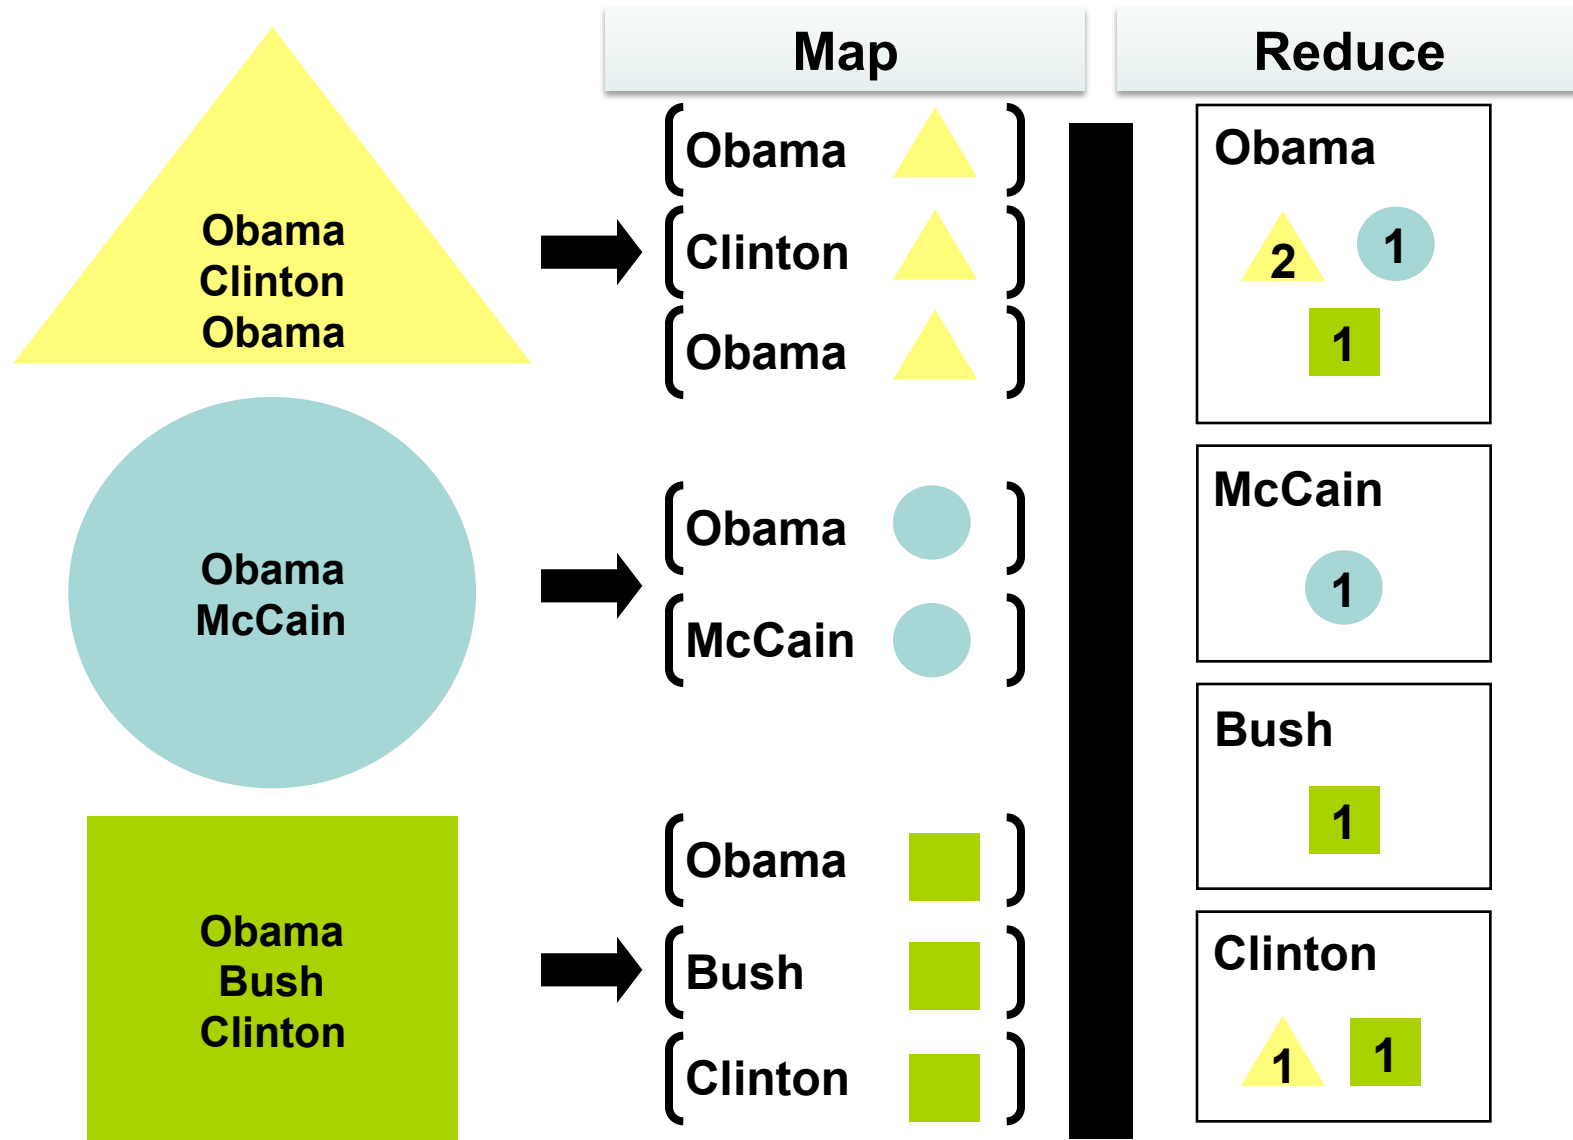
$$\text{sim}(d_i, d_j) = \sum_{t \in V} w_{t, d_i} w_{t, d_j}$$

- But, actually...

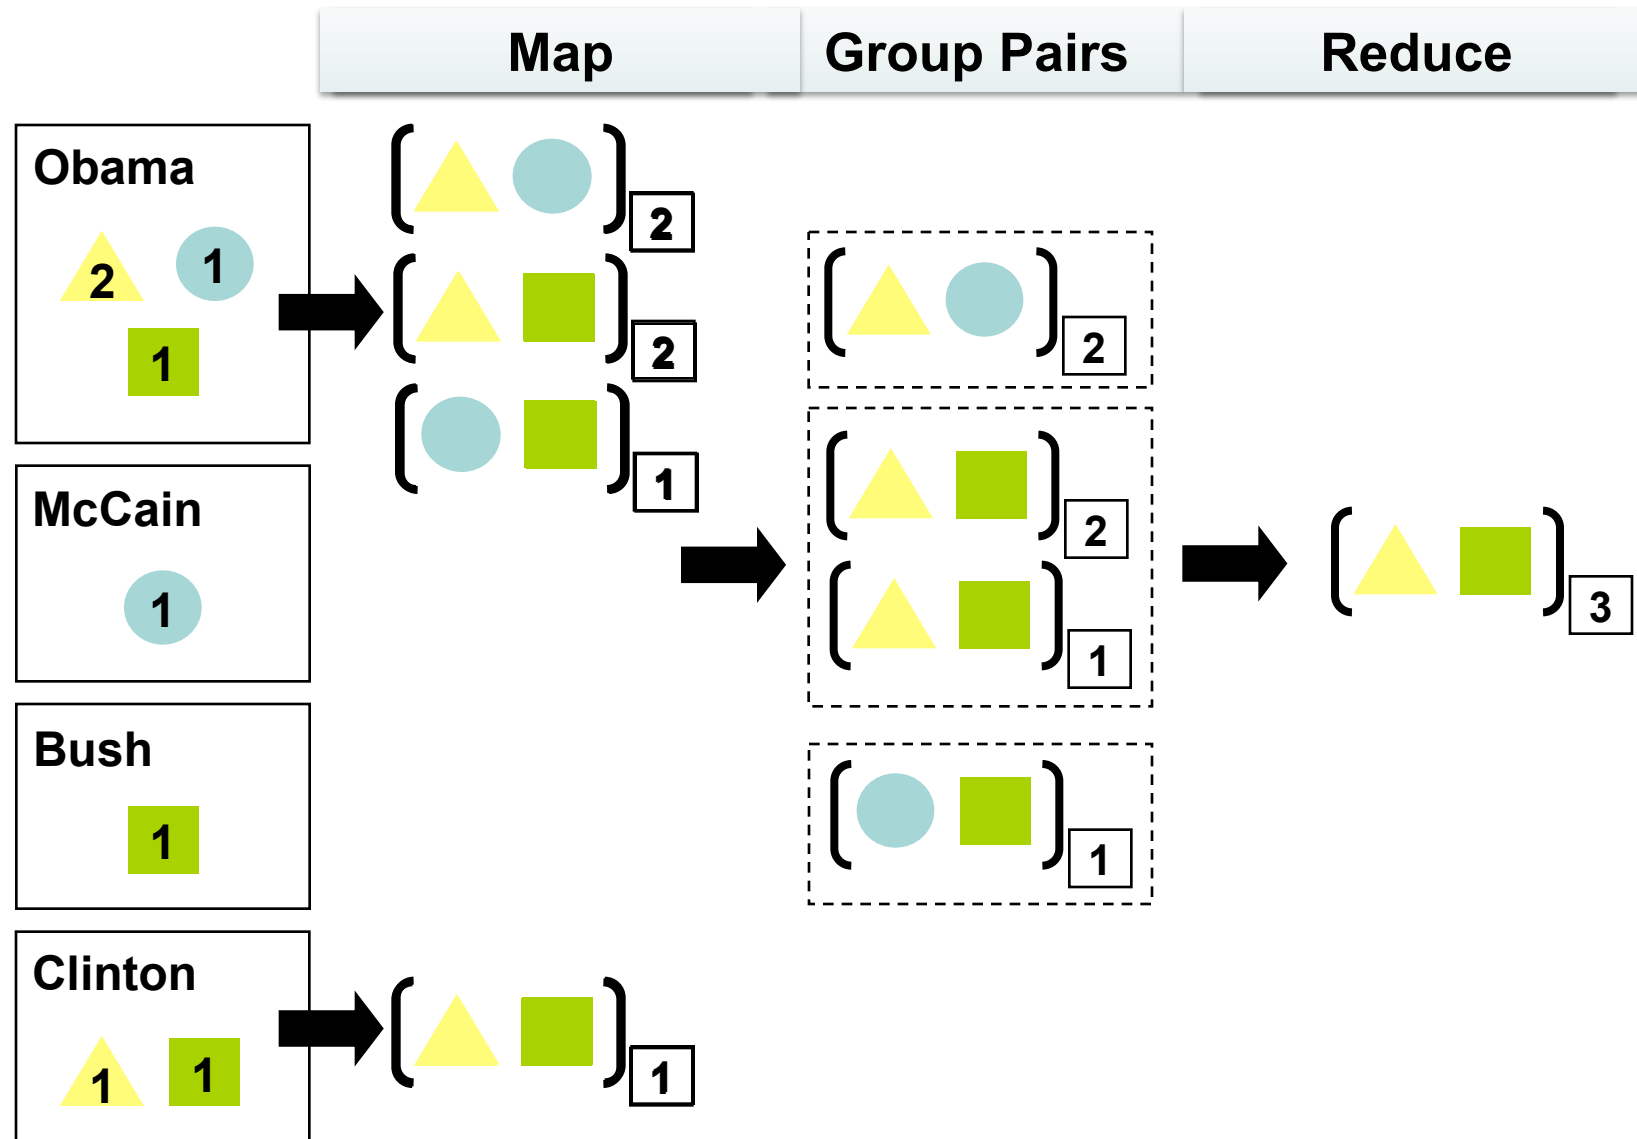
$$\text{sim}(d_i, d_j) = \sum_{t \in d_i \cap d_j} w_{t, d_i} w_{t, d_j}$$

- Two step solution in MapReduce:
 1. Build inverted index
 2. Compute pairwise similarity from postings

Building the Inverted Index




Computing Pairwise Similarity



Analysis

- Main idea: access postings once
 - $O(df^2)$ pairs are generated
 - MapReduce automatically keeps track of partial weights
- Control effectiveness-efficiency tradeoff by dfCut
 - Drop terms with high document frequencies
 - Has large effect since terms follow Zipfian distribution



Data Source: subset of AQUAINT-2 collection of newswire articles;
approximately 900k documents.

Exercise 3: TF-IDF

- Term Frequency
 - Inverse Document Frequency
 - Relevant to text processing
 - Common web analysis algorithm

The Algorithm, Formally

$$\text{tf}_i = \frac{n_i}{\sum_k n_k}$$

$$\text{idf}_i = \log \frac{|D|}{|\{d : t_i \in d\}|}$$

$$\text{tfidf} = \text{tf} \cdot \text{idf}$$

- D – the set of documents
- |D| - the number of documents

Information We Need

- Number of times term X appears in a given document
- Number of terms in each document
- Number of documents X appears in
- Total number of documents

Job 1: Word Frequency in Doc

- Mapper
 - –Input: docname \rightarrow contents
 - –Output: (word, docname) $\rightarrow 1$
- Reducer –Sums counts for word in document
 - –Outputs: (word, docname) $\rightarrow n$
- Combiner is same as Reducer

Job 2: Word Counts For Docs

- Mapper
 - –Input: (word, docname) \rightarrow n
 - –Output: docname \rightarrow (word, n)
- Reducer
 - –Sums frequency of individual n's in same doc
 - –Feeds original data through
 - –Outputs (word, docname) \rightarrow (n, N)

Job 3: Word Frequency in Corpus

- Mapper
 - –Input: ((word, docname), (n, N))
 - –Output: (word, (docname, n, N, 1))
- Reducer
 - –Sums counts for word in corpus
 - –Outputs ((word, docname), (n, N, m))

Thoughts on TF-IDF

- Several small jobs add up to full algorithm
- Lots of code reuse possible
 - Stock classes exist for aggregation, identity
- Jobs 3 and 4 can really be done at once in same reducer, saving a write/read cycle
- Very easy to handle medium-large scale, but must take care to ensure flat memory usage for largest scale

Job 4: Calculate TF-IDF

- Mapper
 - –Input: ((word, docname), (n, N, m))
 - –Assume D is known (or, easy MR to find it)
 - –Output ((word, docname), TF*IDF)
- Reducer
 - –Just the identity function

**Whatever the good tools you have,
algorithms are always the KEY.**

